Case Study:
A Product Experience Scorecard

Making design measurable... and reducing dead clicks by ~30%

Justin Threlkeld ⊙ work from 2024–2025, presented April 2026

The need

Making the impact of design *measurable*

The need

Our director of design wanted to elevate the design practice, telling a the story of design as strategic partner to the business

She wanted to push towards more **impact-based** and **measurable** design work

I was asked to bring actionable *impact data* to the designers
and to help her make it a part of the team's daily work

The need

Okay but like… what exactly are we wanting here?

Ultimately, this request reflected an organizational need to:

1. Quantify the impact of design and product investments
2. Identify areas of user pain and friction
3. Connect product experience metrics to business objectives
4. Align on key performance indicators

The approach

get it started	add meaning	compounding interest	an org-wide signal	world domination?

The approach

get it started	add meaning	compounding interest	an org-wide signal	automate it

The approach

get it started

establish a framework

collect the data

figure out what matters

deliver a metric (maybe two)

Get it started

Where are we starting from?

I asked around, talking to teams about metrics they were tracking, whether they had any key tasks or core workflows they looked at, and what signals might help them better understand how the product was performing.

Get it started

Where are we starting from?

What I learned

- Teams wanted *"a number to measure user experience"*
- Some designers and PMs were familiar with the System Usability Scale
- Most teams didn't have any key tasks or metrics at all

a small wrinkle: missing foundations

Where I would usually reach for key metrics, we had nothing: a wealth of data but no frame in which to use it. No key journeys or user tasks, no core product metrics, no defined jobs to be done.

If my assumptions were off the mark, teams would rightly see the results as irrelevant to their work… and ignore the findings.

a big opportunity: laying foundations

I knew I'd be laying the track while driving the train. I also knew this would be an ongoing effort with outsized impact, provided I could gain the right buy-in.

I embraced the opportunity and outlined an iterative approach: start with something in easy reach, and use it to build momentum, increasing scope, capability, and alignment to org needs each time we ran the study.

Get it started

Finding the data

On the upside, we had tons of data, provided I could get access to it. Amplitude, Fullstory, Salesforce, Snowflake, even Honeycomb itself.

On the downside, the data was in many places, some data had multiple sources, and we had no documented or agreement for which of the ad-hoc definitions for the few things teams were tracking we'd want to use.

Get it started

The qualitative component

From the very start of this project I knew we'd want something more than just product analytics and data exhaust. I'd need a way to reach out to active users.

I needed Intercom access.

Turns out, the team that managed Intercom wanted

✨ *metrics to measure the product experience*

Get it started

Sketching a framework

After the conversations about what we might want to track, how we might want to track it, and how we wanted to organize everything, I suggested we start with the CASTLE framework and add a category for “accessibility”.

It was well suited to B2B SaaS products, accounting for things like cognitive load. learnability, and advanced users, unlike more consumer-focused frameworks.

Seven actionable buckets also felt like the right level of complexity.
Not too many, enough to show the intended scope and future potential.

PX scorecard

Cognitive load

Advanced feature use

Satisfaction

Task efficiency

Learnability

Errors

Accessibility & usability

Get it started

The first version

- UMUX
- Time on Task
- Rage clicks & dead clicks

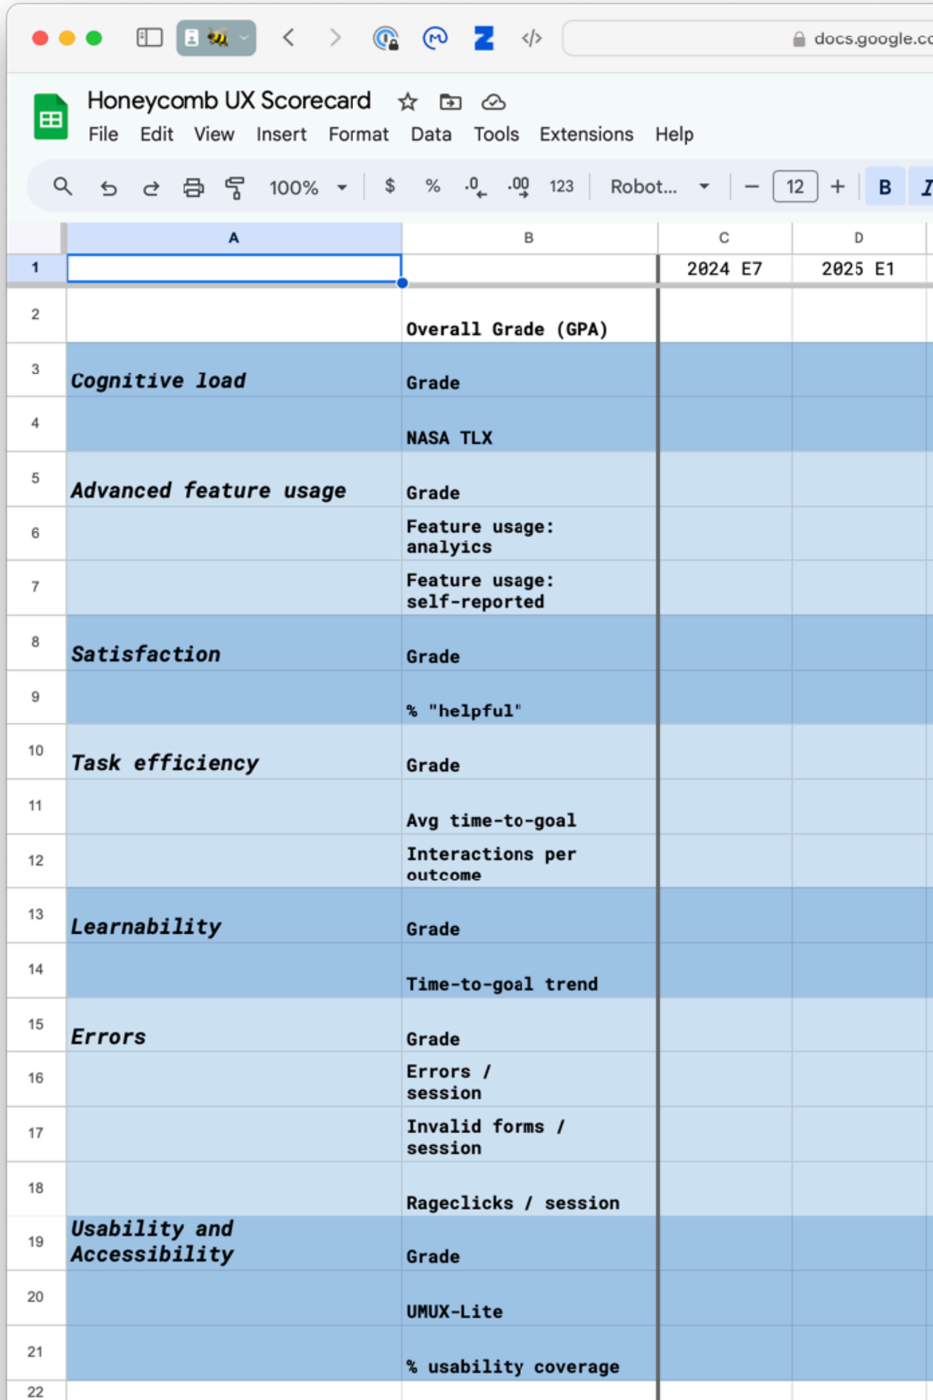

Honeycomb UX Scorecard

	A	B	C	D
1			2024 E7	2025 E1
2		Overall Grade (GPA)		
3	*Cognitive load*	Grade		
4		NASA TLX		
5	*Advanced feature usage*	Grade		
6		Feature usage: analyics		
7		Feature usage: self-reported		
8	*Satisfaction*	Grade		
9		% "helpful"		
10	*Task efficiency*	Grade		
11		Avg time-to-goal		
12		Interactions per outcome		
13	*Learnability*	Grade		
14		Time-to-goal trend		
15	*Errors*	Grade		
16		Errors / session		
17		Invalid forms / session		
18		Rageclicks / session		
19	*Usability and Accessibility*	Grade		
20		UMUX-Lite		
21		% usability coverage		
22				

PX scorecard 2024 Q2

Cognitive load

Advanced feature use

Satisfaction

Task efficiency

Learnability

Errors

Accessibility & usability

The approach

add meaning

expand the metrics

detailed reports with over-time data

activation sessions

The approach

compounding interest

comprehensive metrics

benchmarks (and goals)

key insights & recommendations

PX scorecard 2024 Q2

Cognitive load

Advanced feature use

Satisfaction

Task efficiency

Learnability

Errors

Accessibility & usability

PX scorecard 2024 Q4

Cognitive load

Advanced feature use

Satisfaction

Task efficiency

Learnability

Errors

Accessibility & usability

PX scorecard 2025 Q4

Cognitive load

Advanced feature use

Satisfaction

Task efficiency

Learnability

Errors

Accessibility & usability

Add meaning, build interest

Expand the metrics, make recommendations

- UMUX
- Time on Task++
- Rage clicks & dead clicks
- Product usage errors
- TLX inventory and visual complexity
- Known-pending accessibility issues

Segments by

- user tenure
- account size
- user expertise

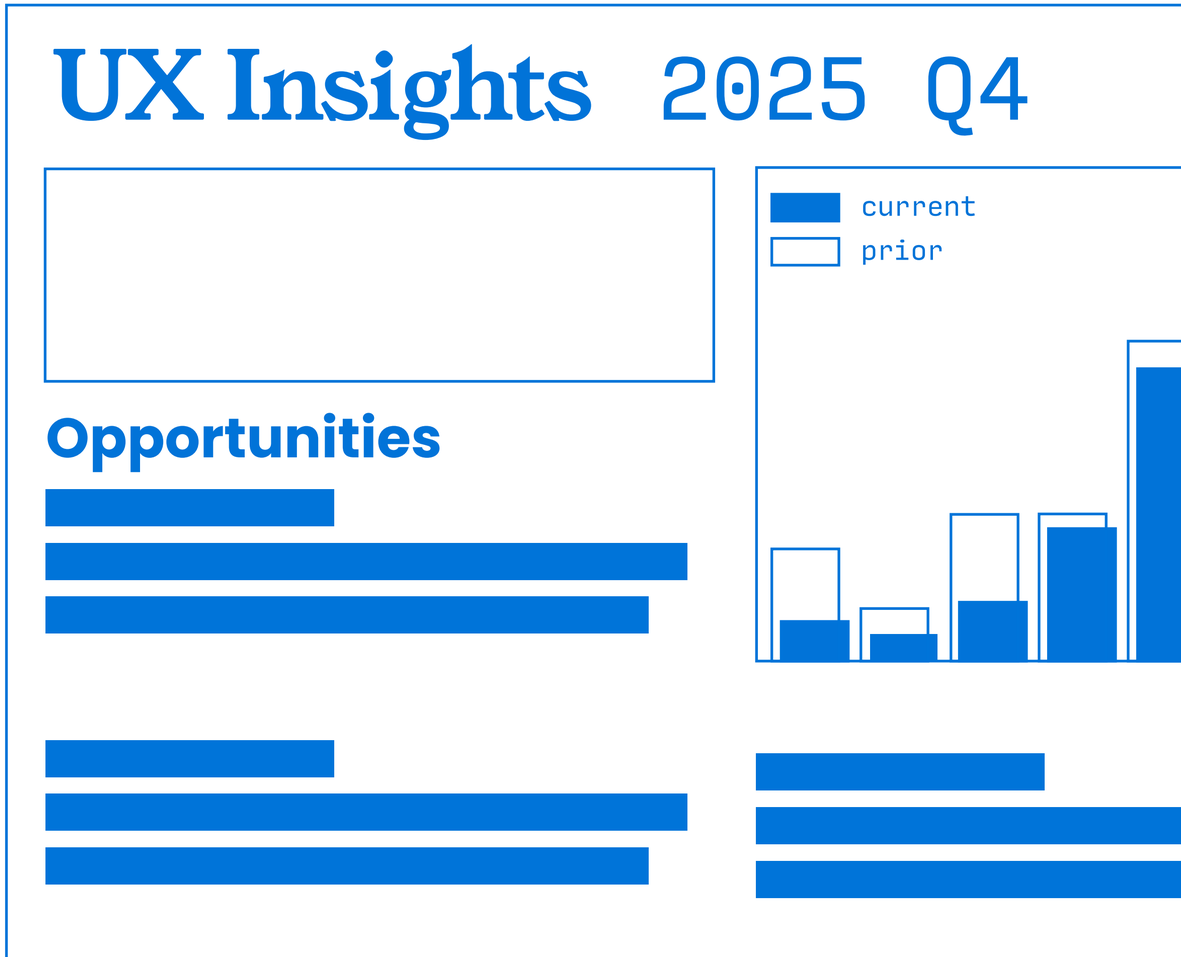

The approach

an org-wide signal

product metrics
in QBRs

metrics for execs

A virtuous cycle

Sharing early to get feedback and interest

Early in the project, I started sharing a topline report in a monthly meeting between product and go-to-market teams. I wanted to

1) generate director and VP level interest for the initiative

2) gather cross functional knowledge that could further contextualize the data

3) get feedback from non-product teams about what insights they might find valuable.

A virtuous cycle

Insight, aligned to need

Gathering early feedback from both executives and product teams, and investing that feedback into each iteration resulted in these readouts quickly becoming aligned to deep organizational priorities with key artifacts shaped to meet needs of leadership, product teams, and cross functional partners like product marketing and customer education.

A metric becomes a KPI

Getting UX represented on the company health dashboard

Based on this work, I was asked to develop of a set of board-facing performance metrics for product teams and initiatives.

The UMUX was a solid fit, as we could easily report it as a single number, break it down into satisfaction and usability components, and segment by user profile data.

We could also leverage a wide pool industry benchmarks for UMUX directly and System Usability Scale equivalents.

Segments and slices

Breaking down the results for different audiences

Enablement teams

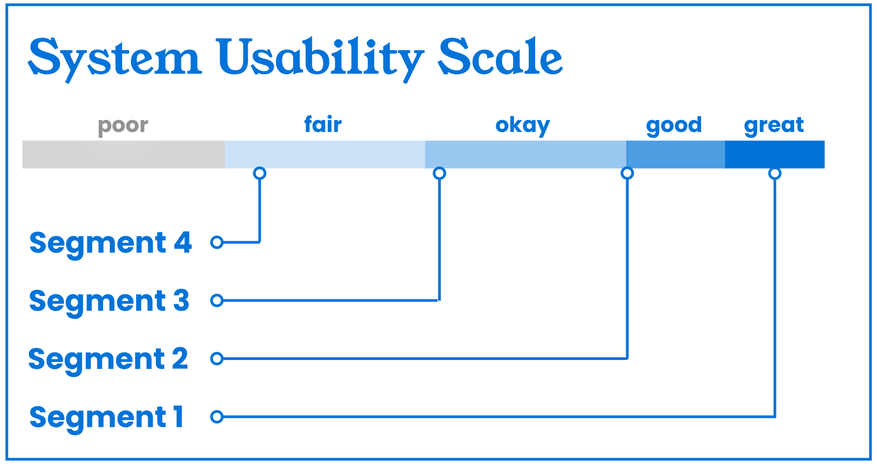

Execs

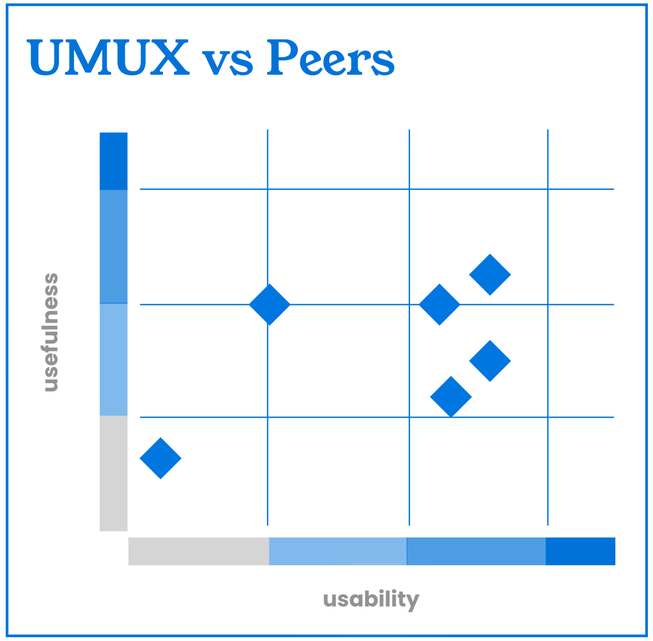

Product teams

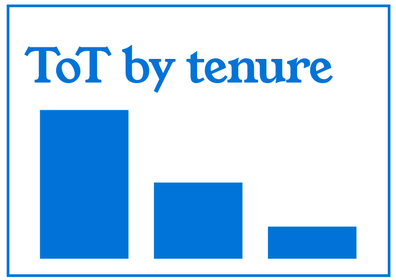

The approach

Insights at a glance, analysis on demand

No more waiting around for the researcher to write the report

Along the way to the full vision of our PX scorecard and insights report, I built out key assets and supporting documentation to allow others to follow along and, ideally, self-serve:

- A Dovetail project containing each scorecard, supporting report, an analysis notes with links to source data
- An Amplitude folder of charts and metrics used for each metric
- An Amplitude dashboard containing the core metrics and links to the full report
- Python notebooks for advanced analysis based on Snowflake data

Putting it on autopilot

Tool consolidation and survey automation

Over the course of the project, Amplitude added session replay, user frustration signals, and surveys to their existing toolset, allowing us to move everything to one tool — allowing me to build a comprehensive dashboard that no longer required copy-pasting numbers between tools.

I could also start running the UMUX survey continuously and piping the results directly to the dashboard.

Having all the data live and in one place meant I could dream a little bigger…

Agentic futures

Experiments on what's next

Elsewhere in the research stack I'd been experimenting with MCP servers and agents, mostly for planning research and searching existing insights.

Having all the data in one place, documentation on how each metric was produced, and a home for analysis notes and summary reporting meant those research planning and insight search tools could take advantage of this data.

It also meant I could start work on proactive, agentic workflows to process the data as it flowed in, injecting insights into Slack conversations and Linear tickets at the speed of the rest of our work.

Impacts

In which I measure the impact of measuring impact.

Impacts

Shaping product priorities, increasing user success

Errors and time-on–task metrics revealed a problematic user journey: about 30% of queries appeared to be run on the wrong dataset.

This aligned to piles of user feedback. And it wasn’t just a frustration issue. Query execution was a large component of our COGS. Even a 1% reduction in this error would have significant impact on our costs per user.

Meanwhile, the AI team was looking for the next workflow to target with LLM augmentation. Based on this data, they picked dataset selection.

Impacts

Shaping product priorities, increasing user success

When I checked back for the next wave...

All users (previous)

28.1%

All users (current)

17.1%

Impacts

Shaping product priorities, increasing user success

This dramatic reduction told us two things

1) The error definition was useful

   a counter-hypothesis was that users might be comparing queries across datasets or have some other reason to quickly change and re-run

2) We were on the right track and should probably re-do the dataset selector

Impacts

Pruning dead clicks

In the first iteration of the scorecard, one metric was an easy target for improvement: dead clicks were very high.

2.24% of all clicks across the product resulted in no action.

This was either a telemetry issue or a problem with our product design, so I shared the findings with our platform and design system team.

It wasn’t the telemetry.

Impacts

Pruning dead clicks

Unfortunately, with the data I had, I couldn't always tell the team which UI components were resulting in dead clicks.

They were already working on a case to prioritize a few design patterns they felt were problematic. but we needed better data to make real progress.

Luckily, they were the right team to fix both sides of this problem.

```
from user_events where type is 'dead_click' select
  page                | element
  page_definition     | element_xpath
count(events)
group by …
```

dead clicks by page

Events	Page	Page Definition
24005	Unknown	Not defined

dead clicks by element

Events	Element	Element xpath
110892	BUTTON	//html/body/div/div[3]/div[2]/div/div/d

Impacts

Pruning dead clicks

While the platform team shipped the general UX improvements they *suspected* would reduce dead clicks, a platform engineer and I worked to improve our Amplitude instrumentation and include component names in user event data.

Armed with data, the team could focus on the patterns they *knew* were underperforming.

We got results.

Dead clicks

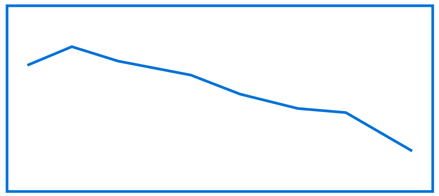

Period	% Dead Clicks
2024 Q2	2.24%
2024 Q4	2.11%
2025 Q2	1.98%
2025 Q4	1.73%

Deep impacts

Teams defining core metrics in QBRs

As metrics became available, teams started citing them in their decision logs and QBRs.

Going further, teams started defining their own metrics, asking for help identifying key user tasks and task flows in usage data.

In the 2024 Q2 iteration, only one team had a user task related UX metric they were tracking. By the 2025 Q4 iteration, 60% of teams had a core UX metric they monitored and reported in QBRs.